

# Mercedes-Benz Greener Manufacturing

The Personal Case Study

## 1. Business Problem



### 1.1 Description

Since the first automobile, the Benz Patent Motor Car in 1886, Mercedes-Benz has stood for important automotive innovations. These include, for example, the passenger safety cell with crumple zone, the airbag and intelligent assistance systems. Mercedes-Benz applies for nearly 2000 patents per year, making the brand the European leader among premium car makers. Daimler's Mercedes-Benz cars are leaders in the premium car industry. With a huge selection of features and options, customers can choose the customized Mercedes-Benz of their dreams.

To ensure the safety and reliability of each and every unique car configuration before they hit the road, Daimler's engineers have developed a robust testing system. But, optimizing the speed of their testing system for so many possible feature combinations is complex and time-consuming without a powerful algorithmic approach. As one of the world's biggest manufacturers of premium cars, safety and efficiency are paramount on Daimler's production lines.

### 1.2 Data Description:

This dataset contains an anonymized set of variables, each representing a custom feature in a Mercedes car. For example, a variable could be 4WD, added air suspension, or a head-up display.

The ground truth is labeled 'y' and represents the time (in seconds) that the car took to pass testing for each variable.

## 1.3 Objective

1. Reduce the time that cars spend on the test bench.
2. To speedier testing, resulting in lower carbon dioxide emissions without reducing Daimler's standards.

## 1.4 Sources/Useful Links

Main Source: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/overview/description>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/overview/description>).

Data: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data> (<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>).

Discussion: <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/discussion>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/discussion>).

## 2. Machine Learning Problem

### 2.1 Data

#### 2.1.1 Data Overview

Get the data from : <https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>  
(<https://www.kaggle.com/c/mercedes-benz-greener-manufacturing/data>)

- Data will be in a file Train.csv and Test.csv
- Size of Train.csv - 3.07MB
- Size of Train.csv - 3.04 MB
- Number of rows and columns in Train.csv = 4210 x 378
- Number of rows and columns in Test.csv = 4210 x 377

#### 2.1.2 Example Data point

[illegible]

- `sklearn.metrics.r2_score(y_true, y_pred, sample_weight=None, multioutput='uniform_average')`
- $R^2$  (coefficient of determination) regression score function.
- Best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of  $y$ , disregarding the input features, would get a  $R^2$  score of 0.0.

## 2.3 Train and Test Construction

We build train and test by randomly splitting in the ratio of 70:30 or 80:20 whatever we choose as we have sufficient points to work with.

# 3. Exploratory Data Analysis

## 3.1 Importing important libraries

```
In [53]: import warnings
warnings.filterwarnings("ignore")
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
import xgboost as xgb
color = sns.color_palette()
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA
from sklearn.metrics import r2_score
%matplotlib inline
from sklearn.svm import SVR, LinearSVC
from sklearn.ensemble import RandomForestRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score

from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score

from sklearn.linear_model import LinearRegression
from sklearn.linear_model import LogisticRegression
import scipy.stats as stats
from sklearn.externals import joblib

from sklearn.pipeline import make_pipeline
from sklearn.linear_model import ElasticNetCV, LassoLarsCV
from sklearn.ensemble import ExtraTreesRegressor, GradientBoostingRegressor
from tpot.builtins import StackingEstimator
```

## 3.1 Reading data and basic stats

```
In [2]: # Load the Drive helper and mount
#from google.colab import drive
# This will prompt for authorization.
#drive.mount('/content/drive')
```

```
In [3]: #Loading data from google drive
#train_df = pd.read_csv('/content/drive/My Drive/Colab Notebooks/Personal Case Study/train.csv')
#print("Train shape : ", train_df.shape)
#test_df = pd.read_csv('/content/drive/My Drive/Colab Notebooks/Personal Case Study/test.csv')
#print("Train shape : ", test_df.shape)
```

```
In [2]: #Loading data from HDD
train_df = pd.read_csv('train.csv')
print("Train shape : ", train_df.shape)
test_df = pd.read_csv('test.csv')
print("Train shape : ", test_df.shape)
```

```
Train shape : (4209, 378)
Train shape : (4209, 377)
```

```
In [3]: train_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4209 entries, 0 to 4208
Columns: 378 entries, ID to X385
dtypes: float64(1), int64(369), object(8)
memory usage: 12.1+ MB
```

```
In [4]: train_df.head()
```

```
Out[4]:
```

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380	X382	X385
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0	1	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0	0	0

5 rows × 378 columns



### Target Variable:

- "y" is the variable we need to predict. So let us do some analysis on this variable first.
- Variable y is of type float
- X0,X1,X2,X3,X4,X5,X6,X8 are of type object
- Rest of the columns are int type
- We will convert [X0,X1,X2,X3,X4,X5,X6,X8] to categorical types and plot to see the distribution of values.

## 3.2 Checking for missing values

```
In [6]: def check_missing_values(df):

    if df.isnull().any().any():
        print("There are missing values in the data")
    else:
        print("There are no missing values in the data")
```

```
In [7]: #calling functions to check missing values on training and test datasets
check_missing_values(train_df)
check_missing_values(test_df)
```


There are no missing values in the data  
There are no missing values in the data

## 3.3 Ploting

### 3.3.1 Ploting y values

```
In [8]: #we are checking 'y' column
plt.figure(figsize=(8,6))
plt.scatter(range(train_df.shape[0]), np.sort(train_df.y.values))
plt.xlabel('index', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.show()
```


""""here we have observed 1 outlier at apporx 260""""



Out[8]: 'here we have observed 1 outlier at apporx 260'

In [9]: *# we again check by visualising in BoxPlot*

```
plt.figure(figsize=(15,5))
sns.boxplot(train_df.loc[:, 'y'])
plt.show()
```



here we have observed 1 outlier


```
In [10]: # we need to remove that outlier
# i have used zscore method and set threshold 10 acc to our data
# https://www.geeksforgeeks.org/scipy-stats-zscore-function-python/

train_df['x'] = np.abs(stats.zscore(train_df.loc[:, 'y']))

outlier_ids = train_df[train_df['x']>10].ID


train_df_final = train_df[~train_df['ID'].isin(list(outlier_ids))]
```

```
In [11]: #now plotting again without outlier
#we are checking 'y' column
plt.figure(figsize=(8,6))
plt.scatter(range(train_df_final.shape[0]), np.sort(train_df_final.y.values))
plt.xlabel('index', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.show()
```



```
In [12]: # we again check by visualising in BoxPlot

plt.figure(figsize=(15,5))
sns.boxplot(train_df_final.loc[:, 'y'])
plt.show()
```




### 3.3.2 Plotting y distribution graph.



```
In [13]: ulimit = 180# we have taken 180 data points
train_df_final['y'].ix[train_df_final['y']>ulimit] = ulimit

plt.figure(figsize=(12,8))#plot size
sns.distplot(train_df_final.y.values, bins=50, kde=True)
plt.xlabel('y value', fontsize=12)
plt.show()
```



```
In [14]: #removing that x helper row for outlier from main row
train_df_final = train_df_final.drop(["x"], axis=1)
```

### 3.3.3 Data type of all the variables present in the dataset.

```
In [15]: dtype_df = train_df_final.dtypes.reset_index()
dtype_df.columns = ["Count", "Column Type"]
dtype_df.groupby("Column Type").aggregate('count').reset_index()
```

Out[15]:

	Column Type	Count
0	int64	369
1	float64	1
2	object	8

Maximum of the columns are integers.  
 8 categorical columns.  
 1 float column (target variable) i.e. 'y'

```
In [16]: #here we can see their types
dtype_df.ix[:15,:]
```

Out[16]:


	Count	Column Type
0	ID	int64
1	y	float64
2	X0	object
3	X1	object
4	X2	object
5	X3	object
6	X4	object
7	X5	object
8	X6	object
9	X8	object
10	X10	int64
11	X11	int64
12	X12	int64
13	X13	int64
14	X14	int64
15	X15	int64

X0 to X8 are the categorical columns


3.4 Plotting these categorical Values

```
In [17]: #https://www.kaggle.com/sudalairajkumar/simple-exploration-notebook-mercedes
var_name = ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']
for val in var_name:
    col_order = np.sort(train_df_final[val].unique()).tolist()
    plt.figure(figsize=(12,6))
    sns.stripplot(x=val, y='y', data=train_df_final, order=col_order)
    plt.xlabel(val, fontsize=12)
    plt.ylabel('y', fontsize=12)
    plt.title("Distribution of y variable with "+val, fontsize=15)
    plt.show()
```


Distribution of y variable with X0




Distribution of y variable with X1




Distribution of y variable with X2




Distribution of y variable with X3





Distribution of y variable with X4



Distribution of y variable with X5





Observation:

- we have observed that X0, X1, X2, X5, X6 and X8 have larger data point.
- X4 and X3 have lesser data point.

## 4. Machine Learning Models

### 4.1 Data preparation

```
In [182]: train_df = pd.read_csv('train.csv')
print("Train shape : ", train_df.shape)
test_df = pd.read_csv('test.csv')
print("Train shape : ", test_df.shape)

y_train = train_df['y'].values
id_test = test_df['ID'].values

usable_columns = list(set(train_df.columns) - set(['ID', 'y']))#taking only important
columns
print(len(usable_columns))

x_train_final = train_df[usable_columns]
x_test_final = test_df[usable_columns]

Train shape : (4209, 378)
Train shape : (4209, 377)
376
```

```
In [183]: # Converting training dataset object categorical values to numerical categorical type
s
#taken help from Link: https://www.kaggle.com/anokas/mercedes-edxgboost-starter-0-5
5

for column in usable_columns:
    cardinality = len(np.unique(x_train_final[column]))

    if cardinality == 1:
        x_train_final.drop(column, axis=1) # Column with only one value is useless so
we drop it.
        x_test_final.drop(column, axis=1)

    if cardinality > 2: # Column is categorical.
        mapper = lambda x: sum([ord(digit) for digit in x])
        x_train_final[column] = x_train_final[column].apply(mapper)
        x_test_final[column] = x_test_final[column].apply(mapper)
```

```
In [184]: # spilting it into 70:30 ratio
X_train, X_test, y_train, y_test = train_test_split(x_train_final, y_train, test_size
=0.3, random_state=42)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(2946, 376)
(1263, 376)
(2946,)
(1263,)
```

## XGBoost

```
In [185]: #taken help from kaggle discussion and kernels for xgboost
#setting up xtrain and xtrain

y_mean = y_train.mean()

d_train = xgb.DMatrix(X_train, label=y_train)
d_cvalid = xgb.DMatrix(X_test, label=y_test)
d_test = xgb.DMatrix(x_test_final)#x_test_final
```

```

In [187]: %%time
#xgb parameters
#just cross validation our model

params = {
    'eta': 0.005,
    'learning_rate': 0.008,
    'max_depth': 4,
    'subsample': 0.9,
    'objective': 'reg:linear',
    'n_estimators': 687,
    'eval_metric': 'rmse',
    'base_score': y_mean, # base prediction = mean(target)
    'silent': 1
}

num_boost_round=2000

#Cross Validation of XGBoost
cv_result = xgb.cv(params,
                    d_train,
                    num_boost_round,
                    nfold = 3,
                    early_stopping_rounds=50,
                    feval=r2_score_metric,#here we have used our metric method
                    verbose_eval=100,
                    show_stdv=False
                    )

```

```

[0]    train-r2:0.00950233    train-rmse:12.3021    test-r2:0.00846467    test
-rmse:12.3007
[100]  train-r2:0.502522    train-rmse:8.71869    test-r2:0.478873    test
-rmse:8.9194
[200]  train-r2:0.61252    train-rmse:7.69449    test-r2:0.571961    test
-rmse:8.08328
[300]  train-r2:0.647986    train-rmse:7.33373    test-r2:0.588661    test
-rmse:7.92327
[400]  train-r2:0.66697    train-rmse:7.13309    test-r2:0.591183    test
-rmse:7.89845
Wall time: 36.1 s

```

**best cv\_result is: train-r2:0.66697 test-r2:0.591183**



```
In [97]: %%time
#Training the model
#taken help from link: https://www.kaggle.com/anokas/mercedes-edxgboost-starter-0-55

#model = joblib.load('model_xgb.pkl')#from load

watchlist = [(d_train, 'train'), (d_cvalid, 'valid')]

model = xgb.train(params, d_train , num_boost_round, watchlist, early_stopping_rounds=60,
                  feval=r2_score_metric, maximize=True, verbose_eval=10)

#joblib.dump(model, 'model_xgb.pkl')#to load
```

[0] train-rmse:12.3026 valid-rmse:13.333 train-r2:0.009819 vali  
d-r2:0.005644  
Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 60 rounds.

[10]	train-rmse:11.7348	valid-rmse:12.8304	train-r2:0.099115	vali
d-r2:0.079199				
[20]	train-rmse:11.2266	valid-rmse:12.3849	train-r2:0.17545	vali
d-r2:0.142034				
[30]	train-rmse:10.7723	valid-rmse:11.9932	train-r2:0.240832	vali
d-r2:0.195436				
[40]	train-rmse:10.3694	valid-rmse:11.6509	train-r2:0.29656	vali
d-r2:0.240714				
[50]	train-rmse:10.0114	valid-rmse:11.3498	train-r2:0.344291	vali
d-r2:0.279445				
[60]	train-rmse:9.69459	valid-rmse:11.0879	train-r2:0.385134	vali
d-r2:0.312319				
[70]	train-rmse:9.41517	valid-rmse:10.8607	train-r2:0.420067	vali
d-r2:0.340208				
[80]	train-rmse:9.16928	valid-rmse:10.6662	train-r2:0.449963	vali
d-r2:0.363631				
[90]	train-rmse:8.95304	valid-rmse:10.4979	train-r2:0.4756	valid-r2:0.3
8356				
[100]	train-rmse:8.76198	valid-rmse:10.3513	train-r2:0.497743	vali
d-r2:0.400651				
[110]	train-rmse:8.59697	valid-rmse:10.2279	train-r2:0.516483	vali
d-r2:0.41486				
[120]	train-rmse:8.45165	valid-rmse:10.1224	train-r2:0.532691	vali
d-r2:0.426872				
[130]	train-rmse:8.32342	valid-rmse:10.0302	train-r2:0.546763	vali
d-r2:0.437258				
[140]	train-rmse:8.21265	valid-rmse:9.95325	train-r2:0.558747	vali
d-r2:0.445863				
[150]	train-rmse:8.11541	valid-rmse:9.88898	train-r2:0.569134	vali
d-r2:0.452995				
[160]	train-rmse:8.03097	valid-rmse:9.83332	train-r2:0.578054	vali
d-r2:0.459136				
[170]	train-rmse:7.95615	valid-rmse:9.78744	train-r2:0.585879	vali
d-r2:0.464171				
[180]	train-rmse:7.89008	valid-rmse:9.74714	train-r2:0.592729	vali
d-r2:0.468575				
[190]	train-rmse:7.83219	valid-rmse:9.71435	train-r2:0.598683	vali
d-r2:0.472144				
[200]	train-rmse:7.77968	valid-rmse:9.68679	train-r2:0.604046	vali
d-r2:0.475135				
[210]	train-rmse:7.73171	valid-rmse:9.66645	train-r2:0.608913	vali
d-r2:0.477336				
[220]	train-rmse:7.6894	valid-rmse:9.64917	train-r2:0.613183	vali
d-r2:0.479204				
[230]	train-rmse:7.65263	valid-rmse:9.63488	train-r2:0.616873	vali
d-r2:0.480745				
[240]	train-rmse:7.6195	valid-rmse:9.62354	train-r2:0.620183	vali
d-r2:0.481967				
[250]	train-rmse:7.58922	valid-rmse:9.61348	train-r2:0.623196	vali
d-r2:0.483049				
[260]	train-rmse:7.56142	valid-rmse:9.60584	train-r2:0.625952	vali
d-r2:0.48387				
[270]	train-rmse:7.5356	valid-rmse:9.60165	train-r2:0.628502	vali
d-r2:0.484321				
[280]	train-rmse:7.50979	valid-rmse:9.59848	train-r2:0.631042	vali
d-r2:0.484661				
[290]	train-rmse:7.48632	valid-rmse:9.59509	train-r2:0.633344	vali
d-r2:0.485025				
[300]	train-rmse:7.4664	valid-rmse:9.59175	train-r2:0.635293	vali
d-r2:0.485383				

[310]	train-rmse:7.44628	valid-rmse:9.59189	train-r2:0.637256	vali
	d-r2:0.485369			
[320]	train-rmse:7.42812	valid-rmse:9.5908	train-r2:0.639024	vali
	d-r2:0.485486			
[330]	train-rmse:7.41023	valid-rmse:9.59082	train-r2:0.64076	vali
	d-r2:0.485483			
[340]	train-rmse:7.39496	valid-rmse:9.59082	train-r2:0.642239	vali
	d-r2:0.485483			
[350]	train-rmse:7.38103	valid-rmse:9.59077	train-r2:0.643586	vali
	d-r2:0.485488			
[360]	train-rmse:7.36523	valid-rmse:9.59208	train-r2:0.64511	vali
	d-r2:0.485348			
[370]	train-rmse:7.35261	valid-rmse:9.59335	train-r2:0.646325	vali
	d-r2:0.485212			
[380]	train-rmse:7.33914	valid-rmse:9.59447	train-r2:0.64762	vali
	d-r2:0.485092			
Stopping. Best iteration:				
[327]	train-rmse:7.41638	valid-rmse:9.59026	train-r2:0.640164	vali
	d-r2:0.485543			

Wall time: 12.5 s

```
In [93]: # Predict on test

#y_pred = model.predict(d_test)
```

```
In [54]: #creating stacked model
#https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LassoLarsCV.html
#https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingRegressor.html
#here we have used GradientBoostingRegressor form ensemble and cross validation LassoLarsCV.
stack_model= make_pipeline(

    StackingEstimator(estimator=LassoLarsCV(normalize=True)),
    StackingEstimator(estimator=GradientBoostingRegressor( learning_rate=0.001,
                                                            loss="huber",
                                                            max_depth=6,
                                                            max_features=0.6,
                                                            min_samples_leaf=18,
                                                            min_samples_split=16,
                                                            subsample=0.7)),
    LassoLarsCV()
)
```

```
In [161]: %%time
stack_model.fit(X_train,y_train)
```

Wall time: 9.25 s

```
Out[161]: Pipeline(memory=None,
  steps=[('stackingestimator-1', StackingEstimator(estimator=LassoLarsCV(copy_X=True, cv=None, eps=2.220446049250313e-16, fit_intercept=True, max_iter=500, max_n_alphas=1000, n_jobs=1, normalize=True, positive=False, precompute='auto', verbose=False))), ('stackingestimator-2', StackingEsti...x_n_alphas=1000, n_jobs=1, normalize=True, positive=False, precompute='auto', verbose=False))])
```

```
In [175]: #final_results=stack_model.predict(x_test_final)
```

```
In [176]: #final_results=model.predict(d_test)
```

```
In [177]: # Predicting R2SCORE
#from sklearn.metrics import r2_score
#r2_score = r2_score(y_test, final_results)#taking r2score on traing data
##print('r2_score = ',r2_score)
#clf = RandomForestRegressor(n_estimators = 60 ,max_depth=5,oob_score=True)
```

```
In [ ]: #predicting on test dataset

final_results=0.10 * model.predict(d_test) + 0.90 * stack_model.predict(x_test_final)
```

```
In [188]: #exporting final results into csv file
test_df = pd.read_csv('test.csv')
csvfile = pd.DataFrame()
csvfile['ID'] = test_df['ID']
csvfile['y'] = final_results
csvfile.to_csv('xgb_stacked.csv', index=False)
```

**Final Test result given by kaggle is: 0.55566**

The screenshot shows the Kaggle website interface for a competition titled "Mercedes-Benz Greener Manufacturing". The header includes the Kaggle logo, a search bar, and navigation links for Competitions, Datasets, Kernels, Discussion, and Learn. The competition banner features a scenic image of a winding road through a mountainous landscape, with the text "Mercedes-Benz Greener Manufacturing" and the question "Can you cut the time a Mercedes-Benz spends on the test bench?". A prize money of "\$25,000" is displayed. Below the banner, there are tabs for Overview, Data, Kernels, Discussion, Leaderboard, Rules, and Team. The "My Submissions" tab is active, showing a table of recent submissions. The table has columns for Name, Submitted, Wait time, Execution time, and Score. A submission named "xgb\_stacked.csv" is listed with a score of 0.55566 and a status of "Complete". A green progress bar indicates the submission is complete. Below the table, there is a link to "Jump to your position on the leaderboard". At the bottom, there is a section for "Step 1: Upload submission file" and a "Make a submission for Ravi Krishna" button. The system tray at the bottom shows the time as 18:03 and network activity.

Kaggle Search Competitions Datasets Kernels Discussion Learn ...

Featured Prediction Competition

**Mercedes-Benz Greener Manufacturing**

Can you cut the time a Mercedes-Benz spends on the test bench?

\$25,000 Prize Money

Daimler 3,835 teams · 2 years ago

Overview Data Kernels Discussion Leaderboard Rules Team My Submissions **Late Submission**

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
xgb_stacked.csv	just now	1 seconds	0 seconds	0.55566

**Complete**

[Jump to your position on the leaderboard](#)

Make a submission for [Ravi Krishna](#)


Step 1  
Upload submission file

5.07 kB/s 18:03

**Feature Importance**

```
In [76]: #https://www.kaggle.com/satadru5/mercedes-benz-xgb-modeling-lb-score-0-54472
fig, ax = plt.subplots(1, 1, figsize=(8, 13))
xgb.plot_importance(model, max_num_features=50, height=0.5, ax=ax)
```


Out[76]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7f5e09445d30>



## 4.2 PCA - Principal component analysis

```
In [0]: # PCA Implementation
pca = PCA(n_components=2)
pca_data = pca.fit_transform(X_train)
```

```
In [0]: cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots(figsize=(10,7))
points = ax.scatter(pca_data[:,0], pca_data[:,1], c=y_train, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```




Observation: here we can see how PCA visualize all the data point far separated from each other, they are not forming tightly group.

### 4.3 T-SNE (t-distributed Stochastic Neighbor Embedding)

```
In [0]: # TSNE Implementation
model = TSNE(n_components=2, random_state=0, perplexity=30)

tsne_data = model.fit_transform(X_train)
```

```
In [0]: cmap = sns.cubehelix_palette(as_cmap=True)
f, ax = plt.subplots(figsize=(10,7))
points = ax.scatter(tsne_data[:,0], tsne_data[:,1], c=y_train, s=50, cmap=cmap)
f.colorbar(points)
plt.show()
```



Observation: here we can see how TSNE visualize all the data point closely attached from each other, They are well grouped .

## 4.4 K-Nearest Neighbors Regressor

```
In [0]: #KNN implementation
#biulding model

knn = KNeighborsRegressor(n_neighbors=5)#k=5 gives best results

knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

r2_score_knn = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(knn.score(X_train, y_train) *100,2)#taking accuracy

results = {'r2_score':r2_score_knn, 'accuracy':accuracy}
print (results)

{'r2_score': 0.414, 'accuracy': 65.57}
```

## 4.5 Support Vector Regressor

```
In [0]: #SVR implementation
from sklearn.metrics import r2_score
clf = SVR()

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print(results)

{'r2_score': 0.384, 'accuracy': 44.76}
```

## 4.6 Random Forest Regressor

```
In [0]: #RFR implementation
from sklearn.metrics import r2_score
clf = RandomForestRegressor(n_estimators = 60 ,max_depth=5,oob_score=True)

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print (results)

{'r2_score': 0.533, 'accuracy': 65.66}
```

## 4.7 Linear Regression

```
In [143]: #Linear Regression implementation
from sklearn.metrics import r2_score
clf = LinearRegression()

clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

r2_score = round(r2_score(y_test, y_pred),3)#taking r2score
accuracy = round(clf.score(X_train, y_train) * 100, 2)

results = {'r2_score':r2_score, 'accuracy':accuracy}
print (results)

{'r2_score': 0.439, 'accuracy': 63.06}
```

## 5. Conclusion



S.no	Model Algo	R2 Score	Accuracy
1.	K-Nearest Neighbors Regressor	0.414	65.57
2.	Support Vector Regressor	0.384	44.76
3.	Random Forest Regressor	0.533	65.66
3.	Linear Regression	0.439	63.06
4.	XGBoost-Stacked	0.55566	-

Here we can see from conclusion that

1. *XGBoost-Stacked perform best in this case by obtaining R2\_SCORE =0.55566 ¶*
2. *Next Our RANDOM FOREST REGRESSOR also perform well with nearly R2\_SCORE=0.533*
3. *After that Knn does better.*
4. *Linear Regression model not suit for this type of problem.*

---XXX---